

Modeling caches in gem5



Before We Start

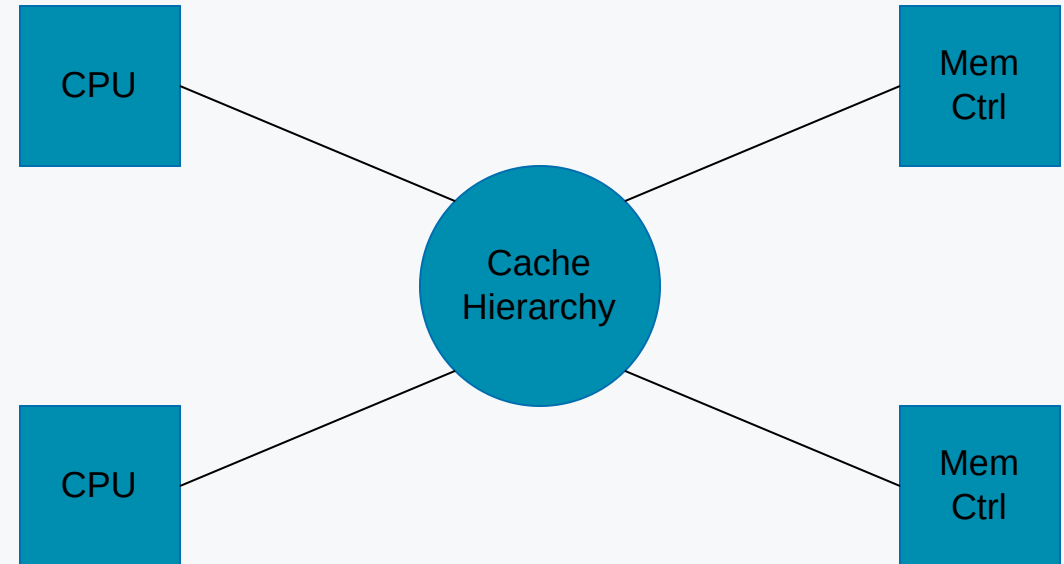
```
cd gem5
```

```
scons build/NULL_MESI_Two_Level/gem5.opt --default=NULL PROTOCOL=MESI_Two_Level SLICC_HTML=True -j17
```



Cache Hierarchy in gem5

1. **Classic Cache:** Simplified, faster, and less flexible
2. **Ruby:** Models cache coherence in detail



Outline

- Background on cache coherency
- Simple Cache
 - Coherency protocol in simple cache
 - How to use simple cache
- Ruby cache
 - Ruby components
 - Example of MESI two level protocol

What is Coherency

A coherence problem can arise if multiple cores have access to multiple copies of a data (e.g., in multiple caches) and at least one access is a write



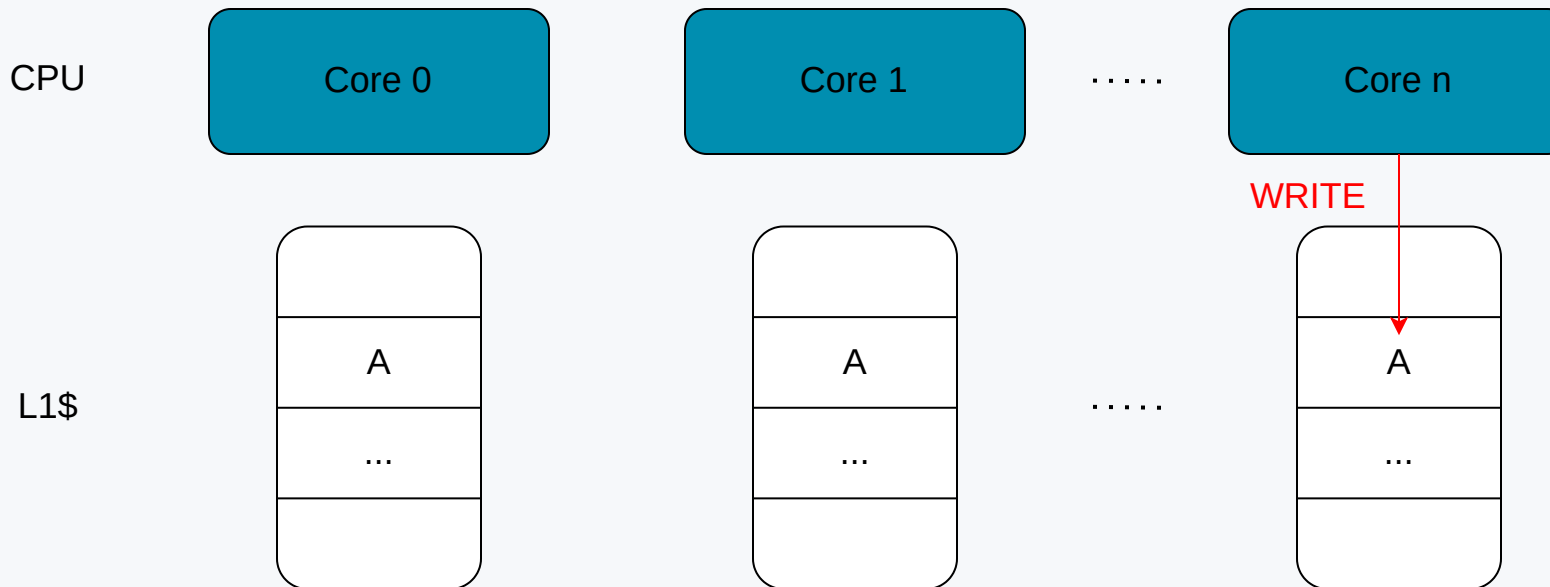
What is Coherency

A coherence problem can arise if multiple cores have access to multiple copies of a data (e.g., in multiple caches) and at least one access is a write



What is Coherency

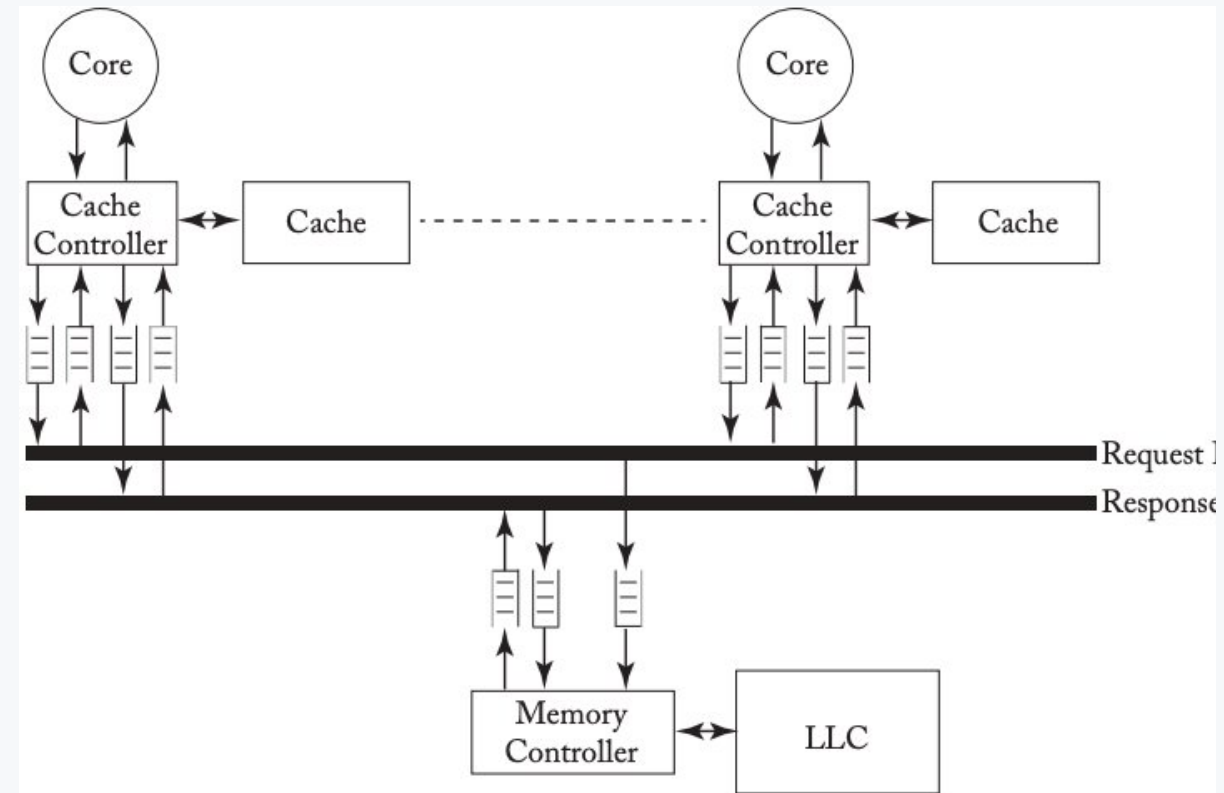
A coherence problem can arise if multiple cores have access to multiple copies of a data (e.g., in multiple caches) and at least one access is a write



- Coherency protocols
 1. Snooping
 2. Directory

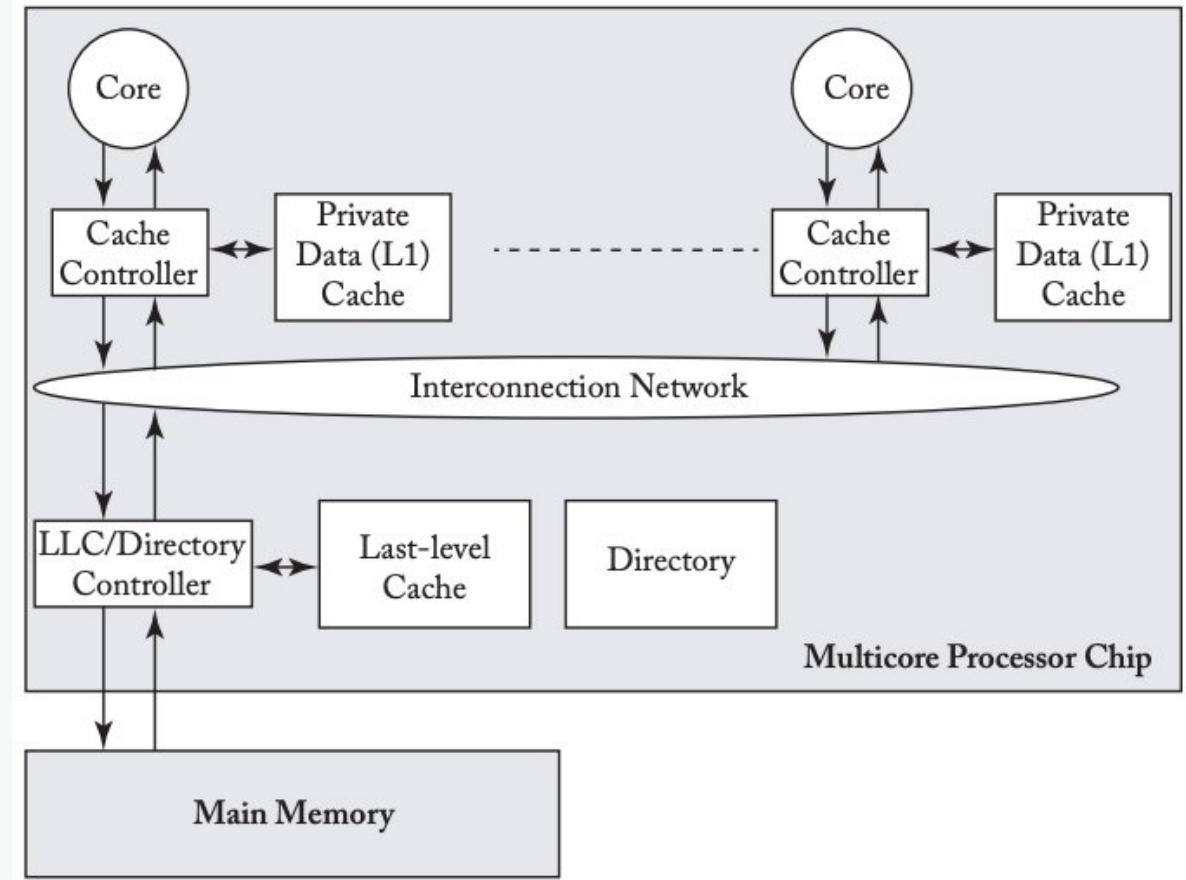
Snoop Protocol

- Each processor snoops the bus to verify whether it has a copy of a requested cache line
- Before a processor writes data, other processor cache copies must be invalidated
- The coherence requests typically travel on an ordered broadcast network such as a bus
- **This technique does not scale since it requires an all-to-all broadcast**



Directory Protocol

- Directory tracks which processor have data when in the shared state
 - Local node where a request originates (interact with CPU cache)
 - Home node where the memory location of an address resides
 - Remote node has a copy of a cache block whether exclusive or shared (interact with CPU cache)
- A general interconnection network allows processor to communicate

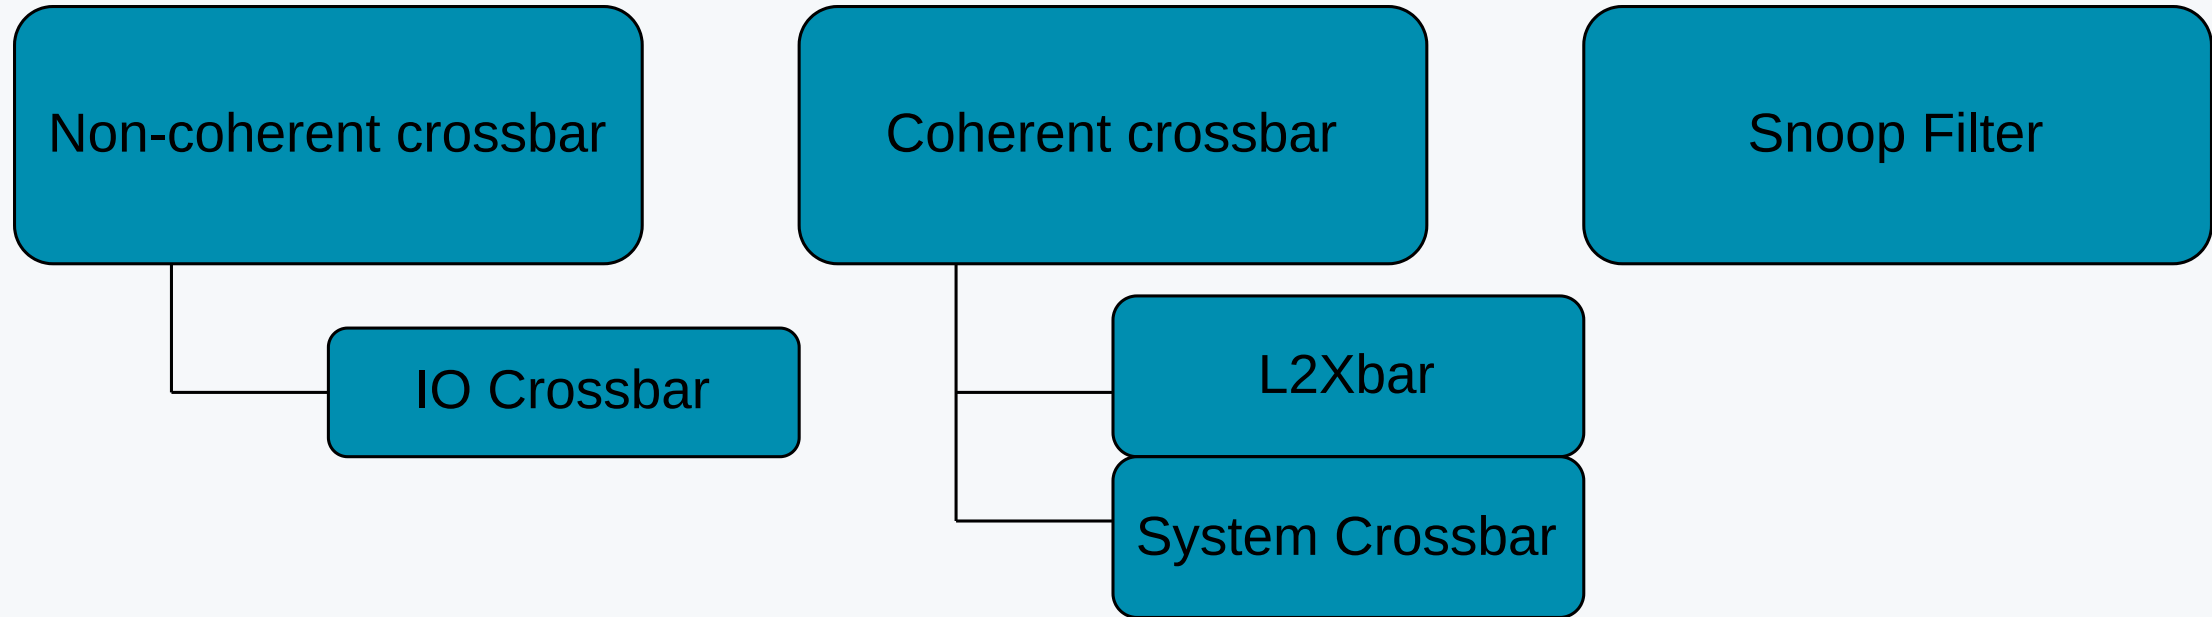




Simple Cache

Snooping Based

Classic Cache: Coherence protocol (Snooping)



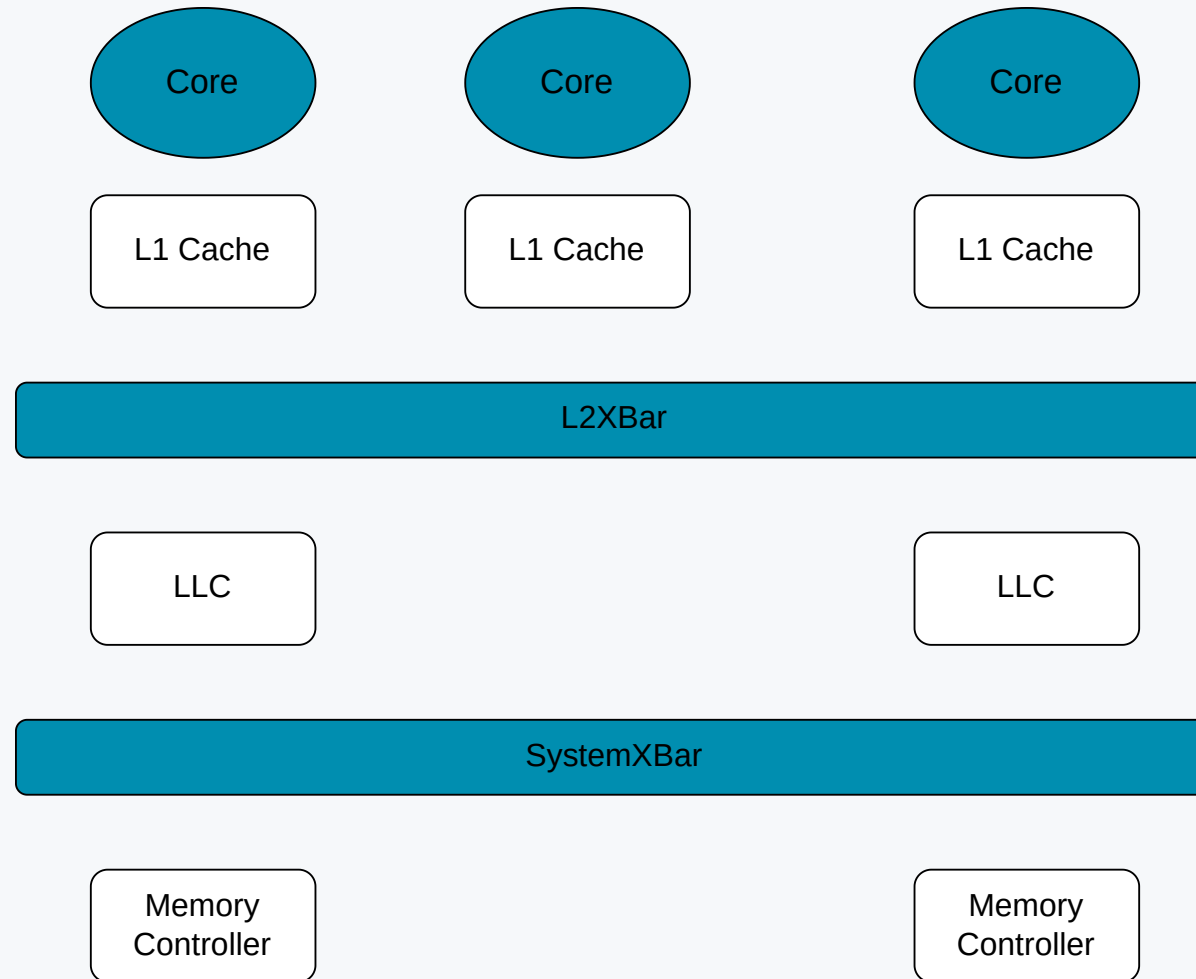
Classic Cache: Coherent Crossbar

- Has snooping request and response bus
- Each core uses the snooping bus to fetch or invalidate a cache line

Classic Cache: Snoop Filter

- Instead of using a snooping bus to find a cache line each Private cache has a snooping directory
- It keeps track of which connected port has a particular line of data
- Instead of snooping the caches it snoops the directory

Example of system with simple cache



Classic Cache: Parameters

- src/mem/cache/Cache.py
 - src/mem/cache/cache.cc
 - src/mem/cache/noncoherent_cache.cc

Parameters:

- size
- associativity
- number of miss status handler register (MSHR) entries
- prefetcher
- replacement policy



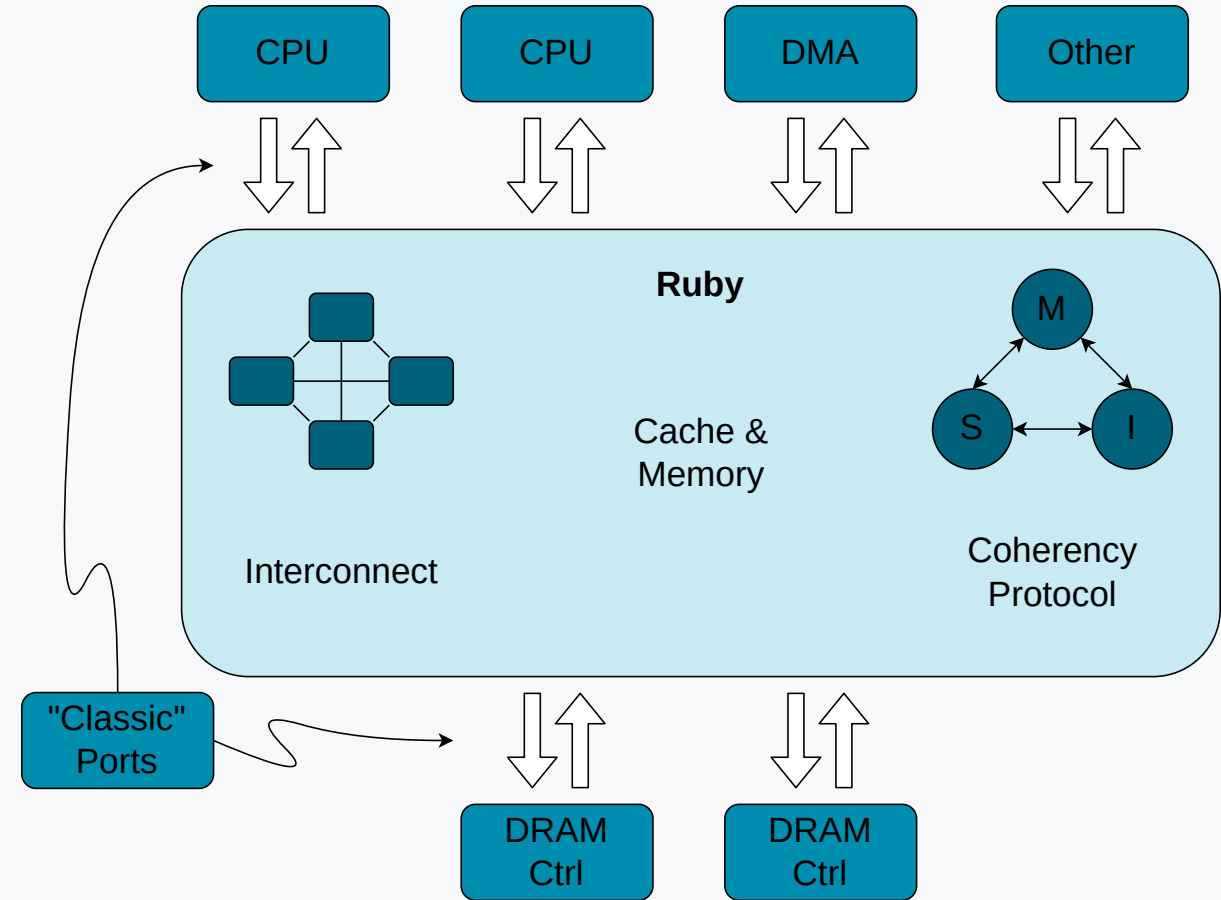
Ruby

Directory Based

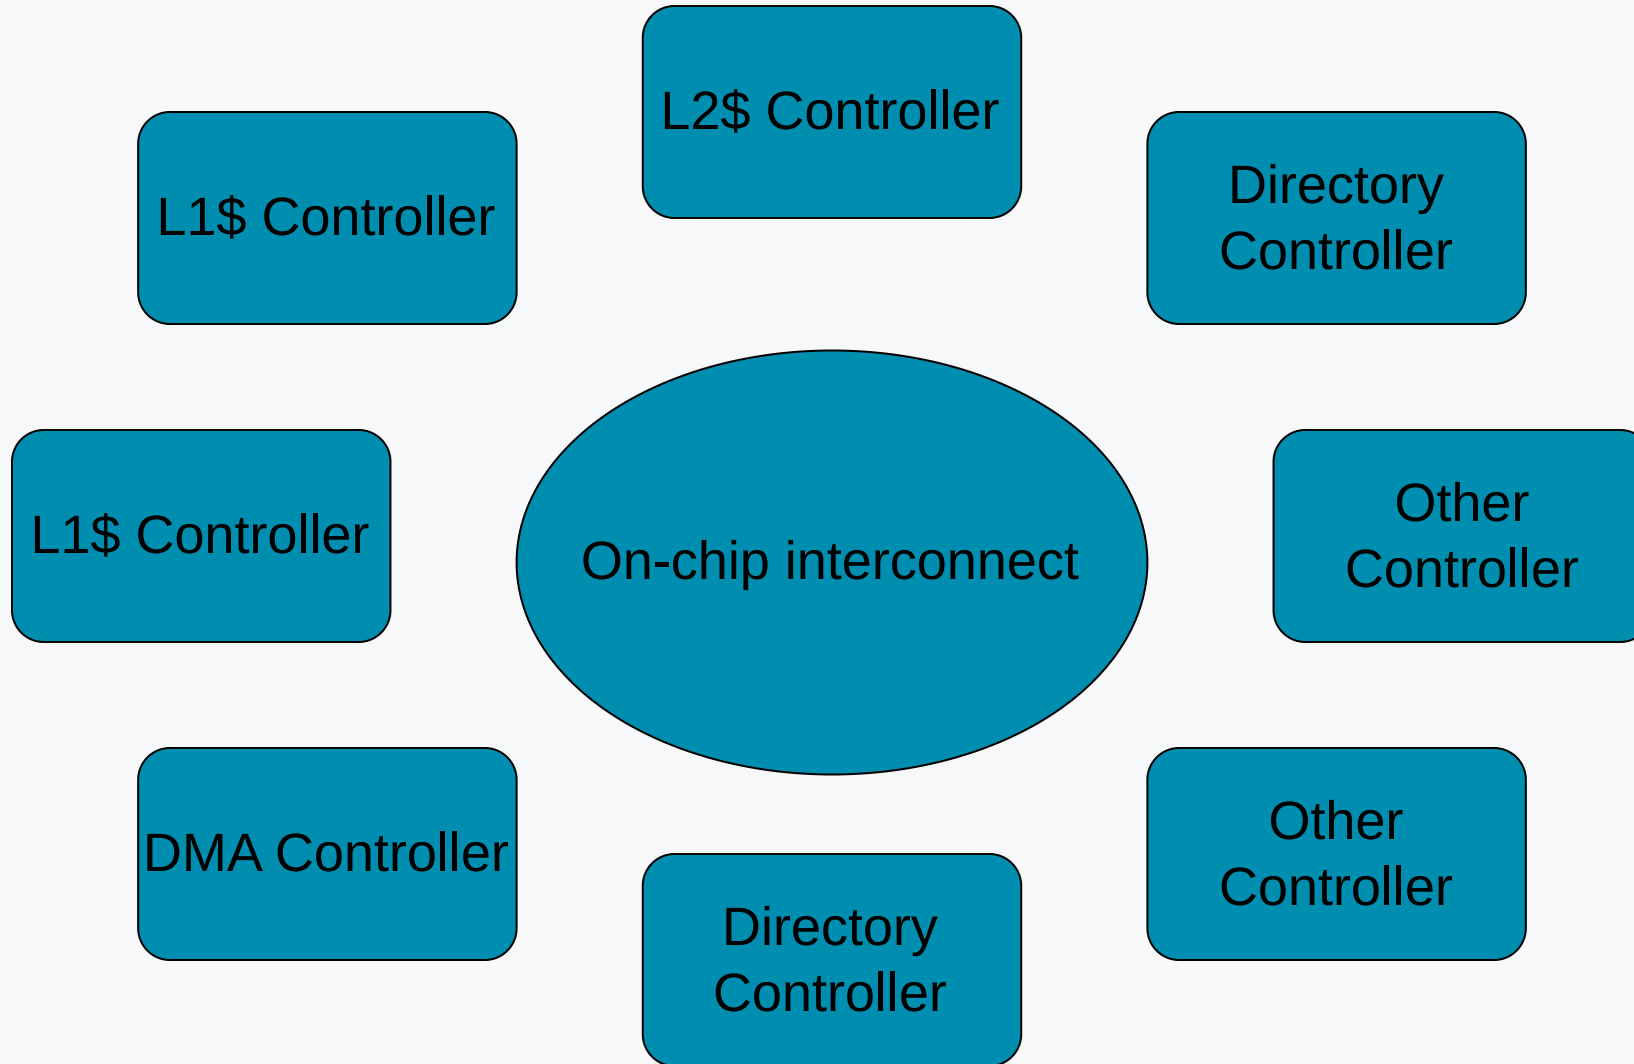


Ruby Cache

1. Coherence Controller
2. Caches + Interface
3. Interconnect



Ruby



Ruby Components

- **Controller models** (cache controller, directory controller)
- **Controller topology** (Mesh, all-to-all, etc.)
- **Network models**
- **Interface** (classic ports)

Ruby Cache: Controller Models

Code for controllers is "generated" via SLICC compilers

Ruby Cache: Example of Controller

Ruby Cache: Caches + Memory



Ruby Cache: Caches + Memory



Ruby Cache: Caches + Memory



Ruby Cache: Caches + CPU

Ruby Cache: Caches + CPU

Ruby Cache System

How to use Ruby

1. Create controllers
2. Create sequencers
3. Connect L1 controllers to sequencers
4. Connect Sequencers to CPUs
5. Connect directories to memory controllers

Example

- Ruby - MESI Two level coherency protocol
- Private L1 cache
- 4 CPUs, 4 private L1 caches
- 1 Shared L2 cache
- 1 Memory channel

Build:

```
cd gem5  
scons build/NULL_MESI_Two_Level/gem5.opt --default=NULL PROTOCOL=MESI_Two_Level -j17
```

Run:

```
cd ../  
.gem5/build/NULL_MESI_Two_Level/gem5.opt materials/using-gem5/04-cache-models/simple_cache_run.py 2 MESITwoLevel 512MB
```

